

Datacenter Extension

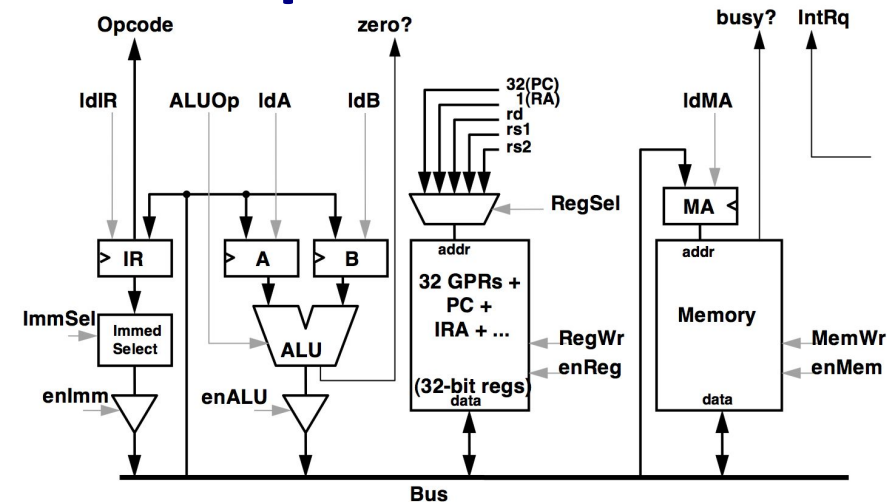
- Move lots of data
- Thousands of threads
- Cache isolation
- Measurement: focus on elapsed time, not CPU time

Dick Sites, Google




The Sodor Processor Collection: Teaching Computer Architecture, Chisel, and RISC-V with Concrete Examples

- **Goals**
 - use Sodor in undergraduate-level computer architecture class (CS152)
 - introduce users to Chisel, RISC-V
- **Menagerie of RV32I processors**
 - 1-stage, 2-stage, 5-stage
 - i. (designed to match lecture slides)
 - 3-stage with synchronous scratchpad (RV32IS)
 - micro-coded bus-based design
- **Sodor available at:**
 - github.com/ucb-bar/riscv-sodor/wiki
 - riscv.org
 - chisel.eecs.berkeley.edu



Chris Celio



advised by Krste & Dave

The Berkeley Out-of-Order Machine: Implementing and Comparing High-Performance Processor Designs using Chisel



- **Goals**

- Create a realistic baseline, strawman RISC-V OoO processor
- Explore energy-efficiency in OoO designs using real RTL that can be pushed through to layout
- provide an open-source implementation for both research and education

- **Why RTL?**

- grounded in reality
- Most other research uses software simulators and energy models

- **Why Chisel?**

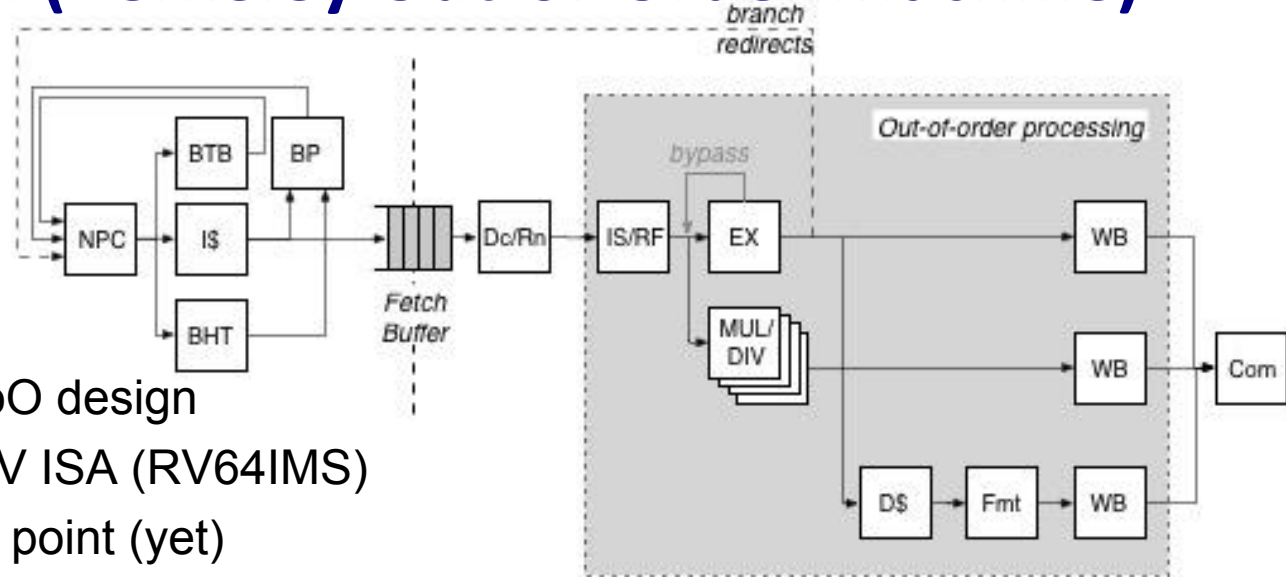
- It's more productive
- more easy to express generators, parameterized hw

Chris Celio



advised
by Krste Asanovic &
Dave Patterson

BOOM (Berkeley Out-of-Order Machine)



- Baseline, strawman OoO design
 - implements RISC-V ISA (RV64IMS)
 - i. no hw floating point (yet)
 - superscalar (parameterizable widths)
 - full branch speculation (BTB/BHT/RAS)
 - load/store queue with store ordering
 - i. loads execute fully OoO wrt stores, other loads
 - ii. store-data forwarded to loads
- > 2.2 GHz @ 45nm (<30 FO4)
- 8,000 lines of Chisel
- fits into existing Rocket-chip stack
 - re-uses significant pieces of Rocket (uncore, caches, frontend, functional units, HTIF, parameter system, FPGA & VLSI flows, and more!)

Status



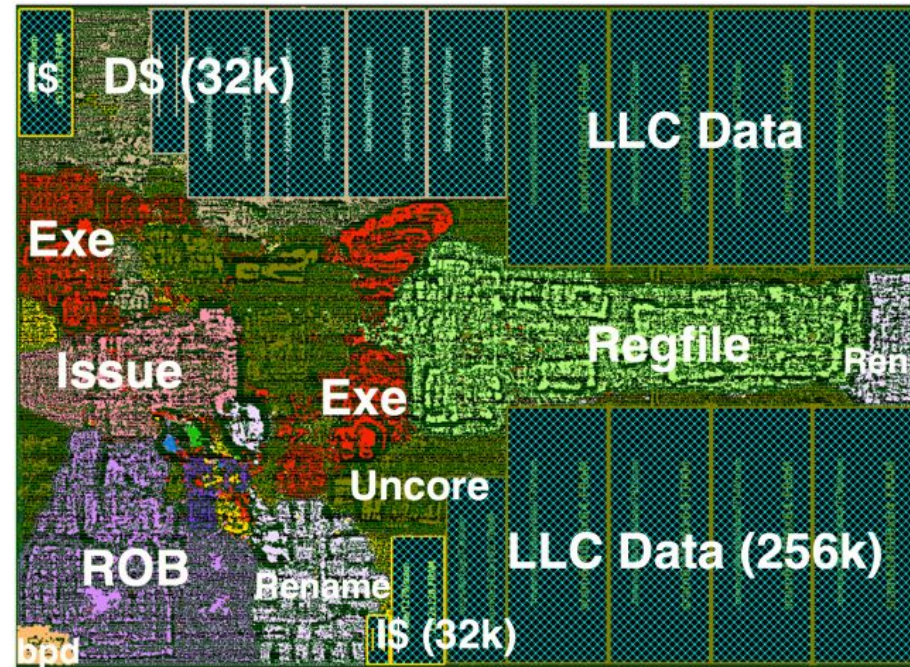
- **Current**

- BOOM has been used in 3 semesters of Berkeley's undergraduate computer architecture class CS152
- Runs on zedboard FPGA (>50 MHz)
- executes dhrystone, Coremark, SPECINT test inputs (~10-100 Billion cycles each)

1.7mm² @ 45nm

- **Next Steps**

- Performance tune
- Finish debugging BOOM
- need story for debugging 100B-1T cycles in
- Explore design trade-offs
- ???
- Paper!




2-wide BOOM layout.

riscv-poky:

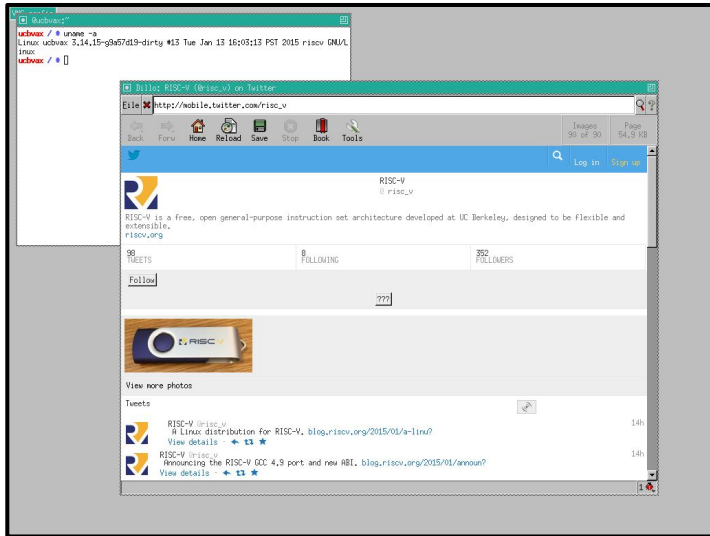
A Linux Distribution for RISC-V

- Compiling an application for RISC-V today:
 - Need to download and compile RISC-V toolchain + Linux
 - Download, patch and cross-compile application and dependencies
 - Create an image to run on QEMU or real hardware
- Problems with this approach:
 - **Error-prone**: Easy to corrupt the FS or get a step wrong
 - **Reproducibility**: Others can't easily reuse your porting work
 - **Rigidity**: If a dependency changes, need to do it all over
- We ported the Yocto Project:
 - A Linux Distribution generator: automate the above steps
 - Part I: **Hundreds of recipes** (scripts that describe how to cross-compile applications for different platforms)
 - Part II: **Build tool** to download, patch, cross-compile and package applications into RPM/DEB, images, SDKs
- Builds hundreds of packages out of the box!




Martin Maas

From riscv-poky to Gentoo Linux



- Why is this cool?
 - We can now easily compile a lot of complex applications, e.g., Python, Perl, LLVM, GCC, etc.
 - Bitbake is a fork of Portage, the Gentoo Linux build tool
 - Since yesterday, we have a working Gentoo port as well!

- If you port a package to RISC-V, **submit a pull request to riscv-poky**. We'll add it to Gentoo as well!
 - Want to grow the list of software available for RISC-V
- Gentoo running a web browser on RISC-V!
→ **Come to the demo!**



Palmer Dabbelt

RICO: Fast Functional Instruction-Set Simulation

Problem Area

- **Fast functional simulation for SW development for many-processor NoC-connected systems**


Current Approach

- **QEMU (Quick EMUlator)**
 - **QEMU design seems resistant to multi-threading**
 - **Host binary translation code generator is arcane**
 - **Code base has lots of PC-related cruft**
 - **Would like to code modules in a newer language**

Our Approach

- **RICO: Code in C and Lua scripting language**
 - **Translate RISC Instructions into Lua**
 - **Let LuaJIT compiler optimize translated code**
 - **Comparable performance to QEMU for 32b RISC**
 - **Multithreaded implementation**
 - **Loadable Lua modules, Lua FFI ifc. to C libraries**


Craig Steele
Tautline LLC



Bluespec demo of BluROCS/Flute on FPGA

- We'll connect remotely to a host+FPGA board in Massachusetts
 - Host: Linux laptop
 - FPGA Board: Xilinx VC707
 - With a PCIe connection
- On FPGA: BluROCS/Flute SoC
- On host:
 - Console I/O
 - Tandem verification with Cissr
 - riscv-gdb:
 - Connecting to BluROCS/Flute
 - Loading an ELF file
 - Setting a breakpoint
 - Running the program (with console I/O)

Rishiyur Nikhil
(with Darius Rad)



FlexPRET: A Predictable Processor for Mixed-Criticality Systems and Precision-Timed IO

Research Area

- Real-time embedded systems, focussing on **timing and safety critical tasks** that interact with sensors and actuators
 - e.g. automotive, avionics


Current Approaches

- RTOS running on traditional embedded processors
 - **unpredictable**: interrupts, unknown processor state
 - **not isolated**: test and verify entire system

Our Approach

- FlexPRET: 5-stage, fine-grained multithreaded RV32I processor
 - **Hard real-time HW threads**: constant rate (configurable)
 - **Soft real-time HW threads**: steal unused cycles
 - Open source, implemented in Chisel, deployed on Xilinx FPGA (Virtex-5, Spartan-6, Zync)

Michael Zimmer



Practical RISC-V Random Test Generation using Constraint Programming

Abstract

A proof-of-concept random test generator for RISC-V ISA is presented. The test generator uses constraint programming for specification of relationships between instructions and operands. Example scenarios to cover basic instruction randomization, data hazards, and non-sharing are presented. The tool integrates the RISC-V instruction set simulator to enable the generation of self-checking tests. The tool is implemented in Python using a freely-available constraint solver library. A summary of problems encountered is provided and next steps are discussed.

Results

- I will be happy to discuss results

Edmond Cote



Objective

- Learn about RISC-V architecture ahead of workshop
- Learn about low-level constraint programming
- Demonstrate proof-of-concept random test generator
- Complete interesting technical project outside of work

Parallel processing Ultra-Low-Power platform

pj/op – mW power envelope, parallel platform for IoT apps

- **Tightly-coupled** clusters of **simple processors** operating in **near threshold**
- Achieve **extreme energy efficiency** on a **wide range of operating points**
- Two chips prototyped in 28nm FD-SOI technology (RVT/LVT)

Davide Rossi



But also...

- **Scalable**: single-cluster (<10 mW) <<>> multi-cluster (10+ GOPS)
- **Programmable**: OpenMP, OpenCL, OpenVX
- **Open**: SW & HW

Prof. Luca Benini


Objective: RISC-V + PULP heterogeneous platform

- *RISC-V based host*
- *PULP many core accelerator (64+ cores, 8+ clusters)*

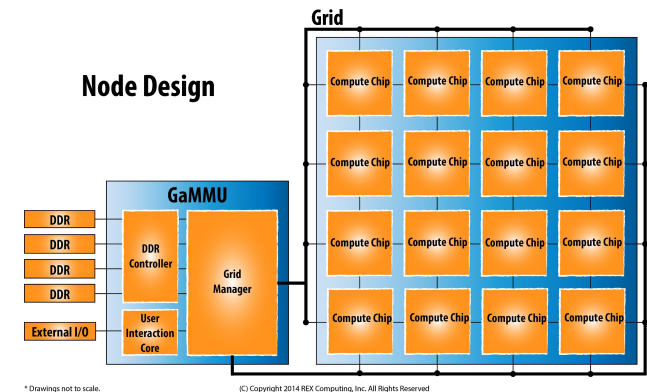
REX Computing Open HPC Processor

<http://rexcomputing.com>

Thomas Sohmers
Paul Sebexen



- HPC Requirements
 - Floating Point and Integer
 - Very high memory bandwidth (For max throughput, 4 words per flop)
 - 50 DP GFLOPs/Watt at system level for Exascale
 - Current state of the art is ~5GFLOPs/Watt
- Processing styles
 - SIMD/Vector doesn't scale for most problems (Amdahl's Law)
 - MIMD/MPMD can (Gustafson's Law... PGAS/SHMEM)
 - Systolic Array + DSP style manycore processing
- Key decisions
 - No (managed) caches: **~30%+ area savings**
 - "You can't fake memory bandwidth that isn't there." --Seymour Cray
 - Global flat memory map for whole chip; can be extended across multiple chips in a node
 - Very simple on-chip mesh network
 - Ideal use would be similar to a Systolic Array, but user is not limited to that.
 - Parallel chip to chip interface
 - Eight 64-bit parallel interfaces per compute chip.
 - Redundant functions are centralized to a separate chip.
- (Hopeful) Results
 - 256 core compute chip (128k scratchpad per core)
 - 256 DP GFLOPs; 85 GFLOPs/Watt chip level
 - 53 GFLOPs/Watt node level
 - 384 GB/s aggregate compute chip bandwidth



Near future:

Full spec: March 2015

Cycle accurate simulator: ~March 2015

Prototype (16 core + interfaces) tapeout ~Q4 2015.

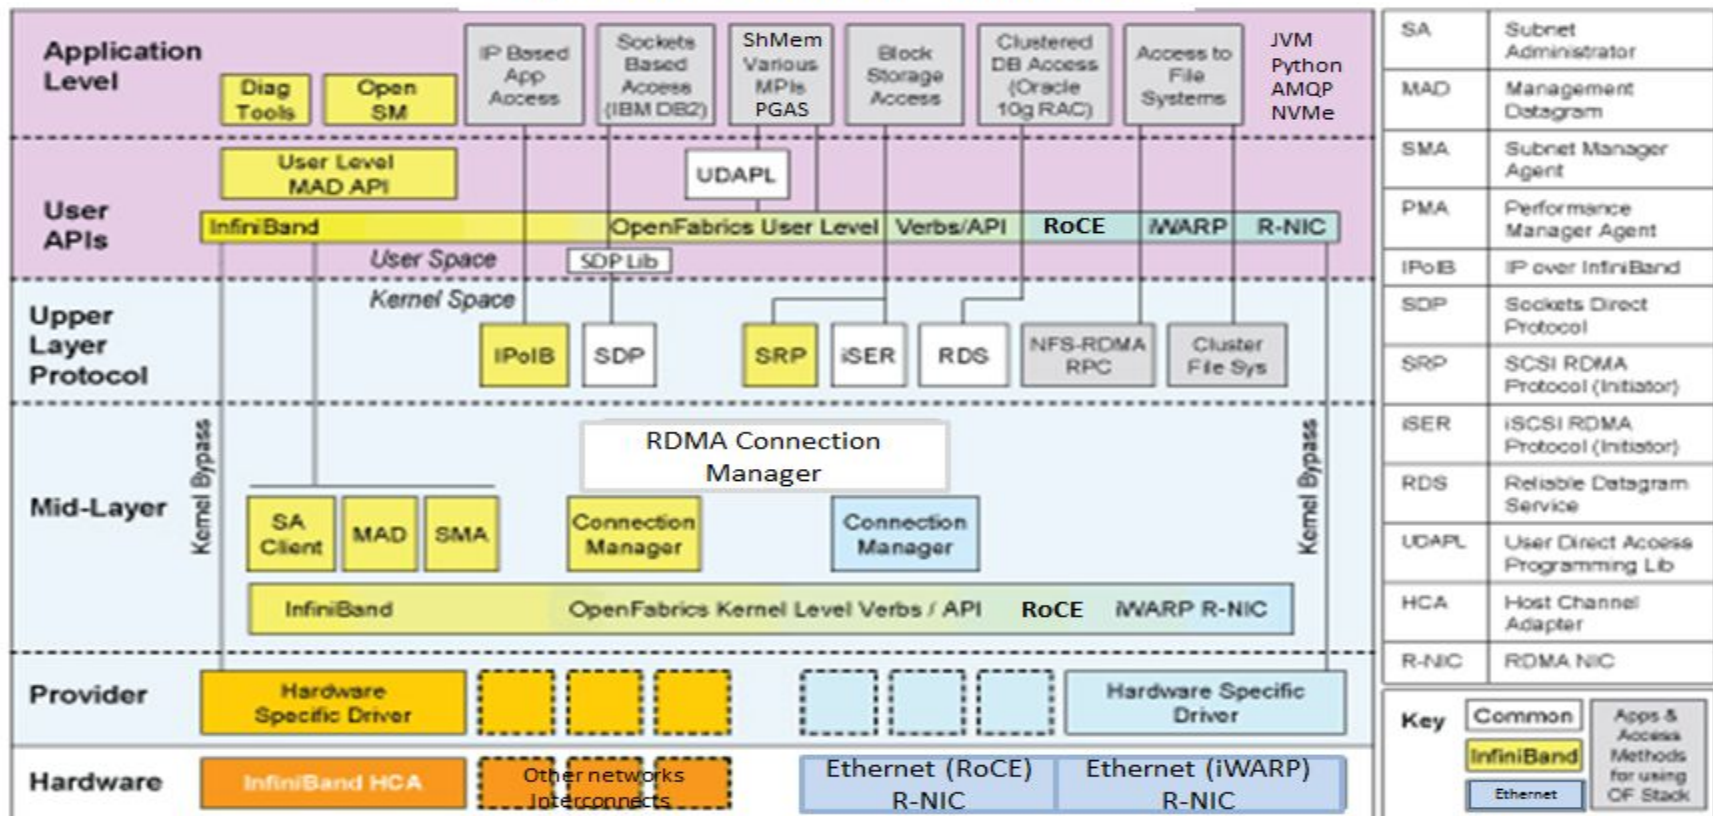


bill.boas@openfabrics.org




Deployed on ~ 50% of Top500 machines

OpenFabrics Software - Linux Stack



Universal Integers

- C int type
 - counting objects in memory ✓
 - arbitrary external quantities ✗
 - favors performance over correctness
- Goal: Add **unbounded integers**
 - Small Integers *as efficiently as C int*
 - Big Integers in software eg GMP
- Why?



David Hossack

Understanding Integer Overflow in C/C++


Will Dietz, Peng Li, John Regehr, Vikram Adve

Proceedings of the 34th International Conference on Software Engineering (ICSE), Zurich, Switzerland, June 2012

- How?
Encoded tagged union { small_int, pointer }
- Student project?


Probabilistic Programming with ISA-Level Primitives

2 new elementary logic gates!
(one that will annihilates the universe)



A probabilistic processor design!
(impossible to physically construct)

**Programming model
based on fictional technology!**
(Infinite Improbability Drive)



**RISC-V + probabilistic instructions
will save the day!**
(ARMv6 sorta worked, but it was
messy)

Adam M. Smith
[@rndmcnly](#)



Probabilistic Programming with ISA-Level Primitives

- **Probabilistic programming is a thing these days**
 - using programs to define probability distributions
 - fancy algorithms answer questions about them
- **PP languages are still out of reach for most programmers**
 - unfamiliar statistics vocabulary
 - weak tool support
 - black-box model: poor search/inference performance
 - white-box model: severe restrictions on computation
- **Dream: black-box flexibility + white-box efficiency**
- **Idea: implement a probabilistic processor**
 - expose new instructions with C wrappers
 - reuse existing languages and compilers to produce machine code
 - apply symbolic simulation to reason over all possible executions
(a twist on probabilistic model checking)
- **Prototypes:**
 - An ARMv6 probabilistic processor (barely works)
 - can generate and solve math textbook exercises given a student's mental model described in C
 - Python library for rapid prototyping of probabilistic software
- **Future:**
 - switch to RISC-V; use a real HDL; apply logic minimization; ...
 - think like a processor designer, not like an AI/PL/stats researcher
 - interesting new uses for an open ISA; fresh ideas for AI/PL/stats


Adam M. Smith
[@rndmcnly](#)



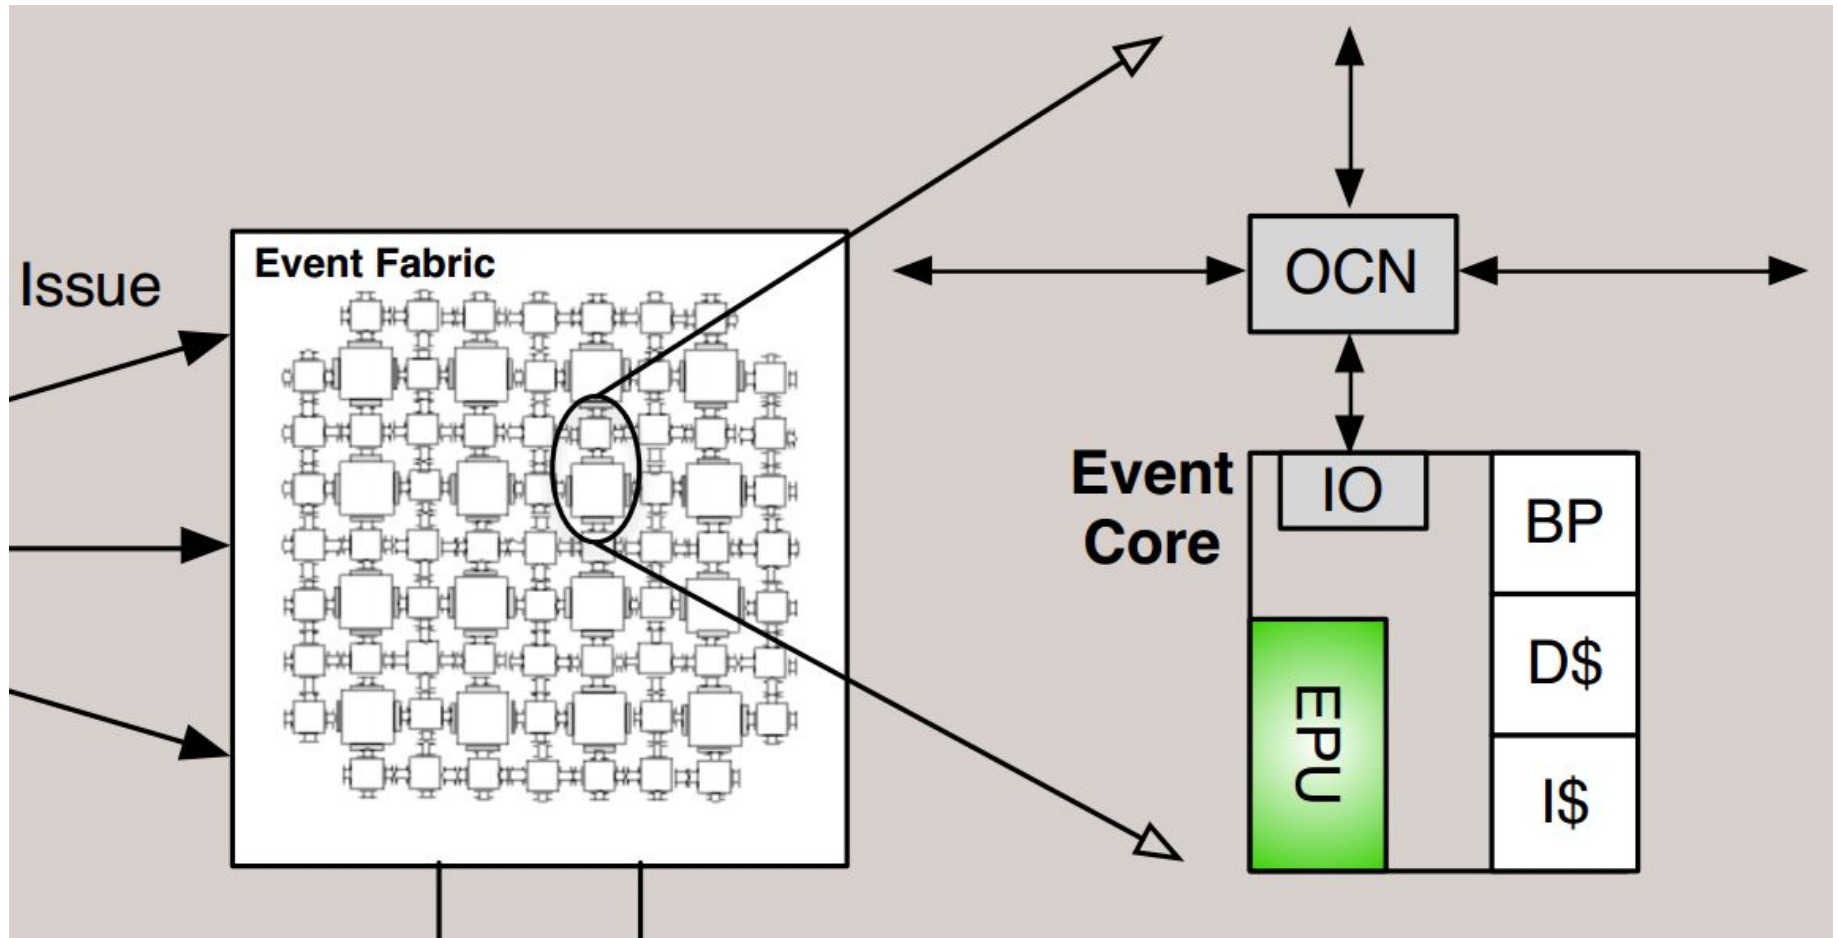
Eve: An Event-Driven Architecture For Asynchronous Programming Models

- Observation:
 - Architecture innovations driven by applications characteristics and software programming model upheaval (e.g., GPUs)
- Event-driven model as a new programming model and application class
 - IoT, sensor-rich environment
 - Mobile devices
 - Cloud (e.g., Node.js)
- General-purpose CPUs are ill-suited
 - Locality
 - Parallelism

Yuhao Zhu
yuhaozhu.com
UT Austin

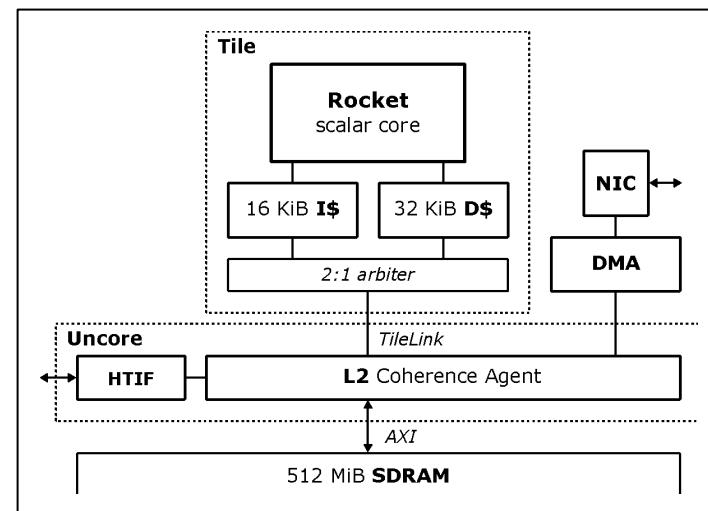


Eve: An Event-Driven Architecture For Asynchronous Programming Models



Hardware Acceleration of Key-Value Stores

- Datacenter apps, path through CPU/kernel/app $\approx 86\%$ of request latency
- Goal: Serve popular Key-Value Store GET requests without CPU
- Soft-managed cache attached to NIC, RoCC CPU interface
- Benchmarking on FPGA:
 - RISC-V Rocket @ 50 MHz
 - NIC from TEMAC/PCS-PMA + 1 Gb SFP
 - Hardware KV-Store Accelerator
 - Traffic Manager, DMA Engine
- Written in Chisel



Base System: First Physical RISC-V System with Networking Support

Sagar Karandikar

Howard Mao

Albert Ou

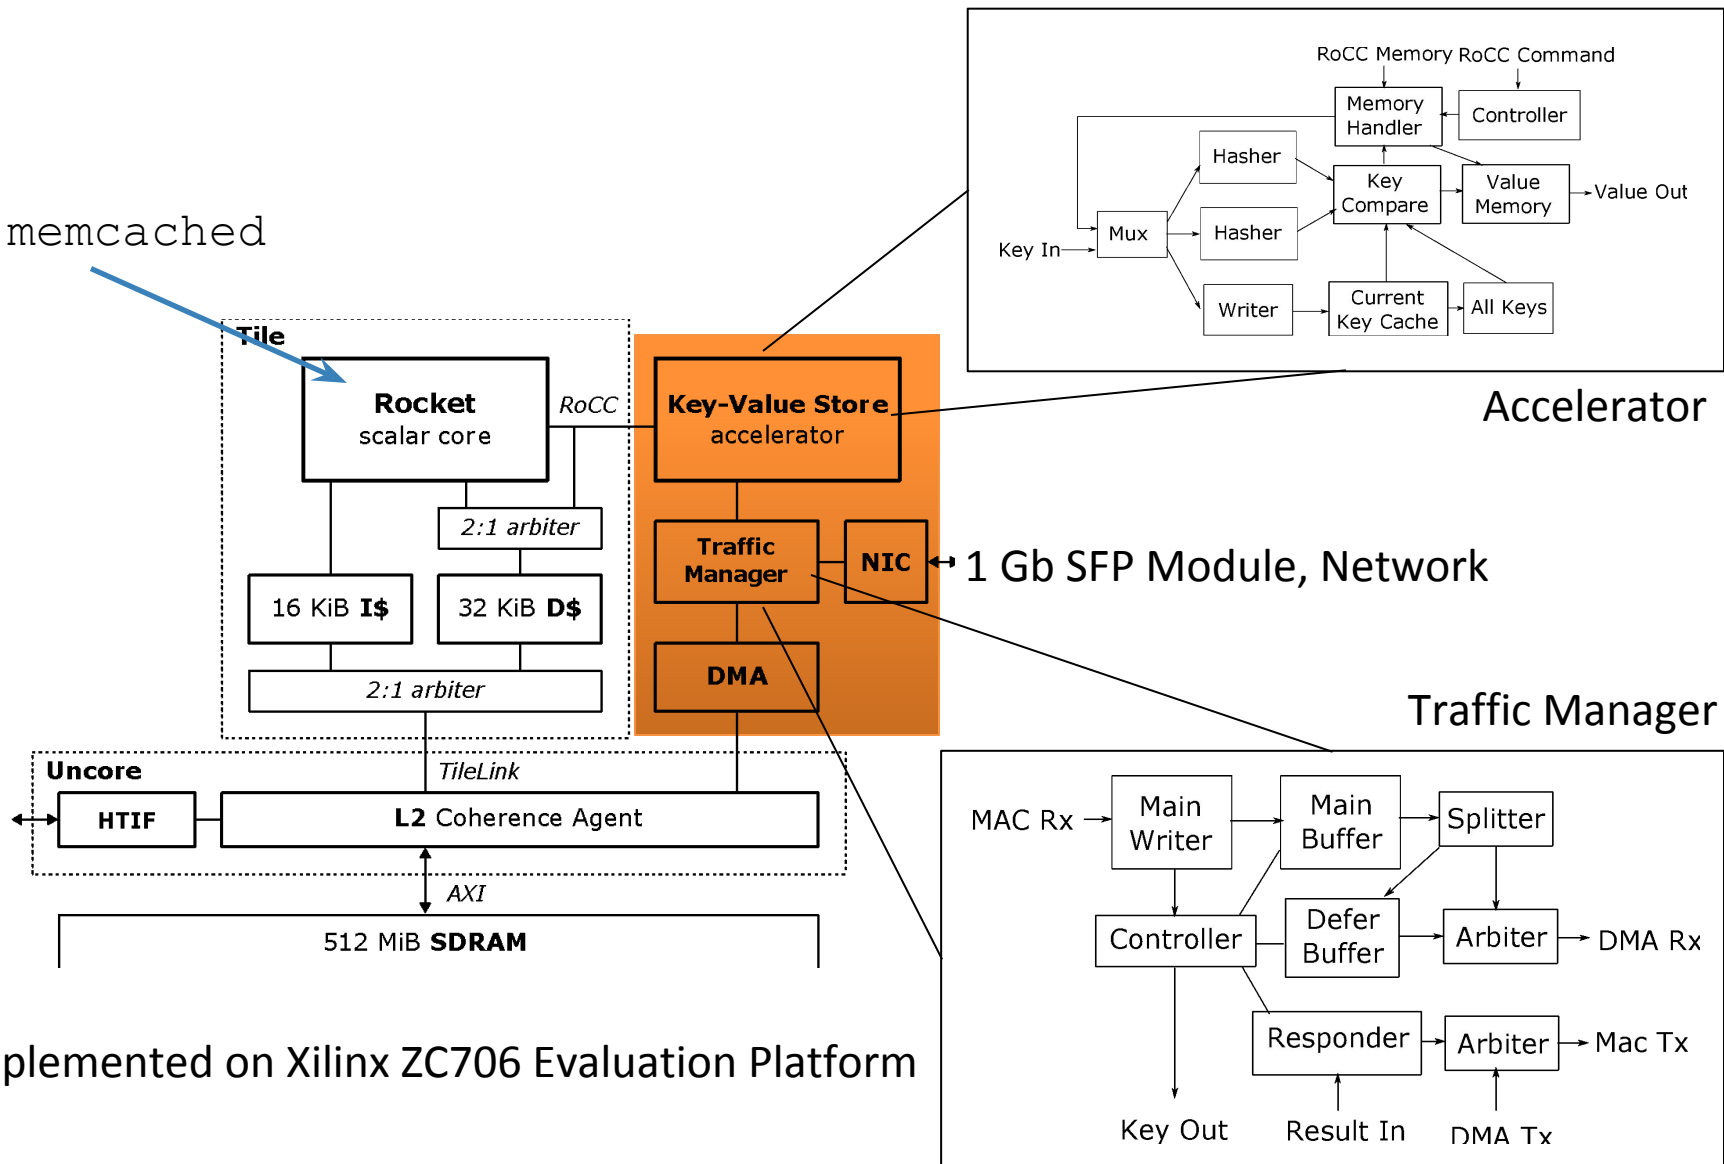
Yunsup Lee

Soumya Basu

ASPIRE Lab, UC Berkeley

Hardware Acceleration of KV Stores - System Design

Runs memcached



Implemented on Xilinx ZC706 Evaluation Platform

Hardware Acceleration of KV Stores - Evaluation

- Raw Performance: $\sim 10\times$ Latency Reduction
 - Software: $\sim 1700\ \mu\text{s}$
 - Accelerator: $\sim 150\ \mu\text{s}$
- Replacement policy is software-tunable
- Realistic benchmark based on appx. of internal Facebook traces:
 - Handled 40% of requests @ 10x latency reduction

